


Deepmac

Utilizing deep Q-learning to create agent playing Pac-Man

- Jakub Skalski
- Paweł Szmergala

What is Q-learning?



$$\underline{Q(S_t, A_t)} \leftarrow \underline{Q(S_t, A_t)} + \alpha [\underline{R_{t+1}} + \gamma \underline{\max_a Q(S_{t+1}, a)} - \underline{Q(S_t, A_t)}]$$

New
Q-value
estimation

Former
Q-value
estimation

Learning
Rate

Immediate
Reward

Discounted Estimate
optimal Q-value
of next state

Former
Q-value
estimation

TD Target


TD Error

Models and data




- Data is generated through gameplay and stored in memory in form of [state, action, reward, next state] **transitions**.
- Model consists of two neural networks both with the same architecture (double Q-Learning).
- Training agent takes action according to the **ϵ -greedy** strategy.
- Input is read from an image and processed before being fed to the convolutional neural network.






Architecture



Successful training process




Unsuccessful training process



Conclusions



Deep reinforcement learning is hard :)



Input format matters way more than expected



Lots of hyperparameters to manage





Ambiguous diagnostics

Results

- Agent shows signs of learning.
- Results are way better than those of a random playing agent.
- We could not make the agent learn from non-image input format of the game which was disappointing.
- Agent looks promising yet training is very computationally expensive.

[Sauce code](#)



Semi one-hot
encoded game state